

Weka and Machine Learning

Abram Hindle

2019-02-12 Fri

Outline

Weka

Machine Learning Classifiers

Learners

Evaluation/Validation

Let's Experiment!

Weka Stuff

Licensing

- ▶ Licensed under Creative Commons CC-BY-SA 3.0
- ▶ Some text and images copyright Wikipedia 2012

Intro

- ▶ AI that learns from data
 - ▶ Learn what spam looks like to filter it out
- ▶ Classify data into types
 - ▶ Learning spam
- ▶ Cluster data by similarity
 - ▶ Finding messages that are similar to spam
- ▶ Find important and distinct properties of the data.
 - ▶ Viagra is a spam keyword!

Kinds of ML

- ▶ Supervised
 - ▶ we give it classified examples and hope it can classify more
- ▶ Unsupervised
 - ▶ labels unknown, let the algorithm find them
- ▶ Semi Supervised learning
 - ▶ labelled and unlabelled.
- ▶ Representation Learning
 - ▶ How to represent data
- ▶ Reinforcement Learning
 - ▶ policies to reward the learner

WEKA

- ▶ ML Toolkit
- ▶ Good for initial exploration if something will work
- ▶ Can use it as a library later
 - ▶ Java
- ▶ Can use it from the commandline
- ▶ Nice way to debug classic ML learners
- ▶ Major Weakness: Tuning

Some Kinds of Learners for Classification

- ▶ Tree Based
 - ▶ C4.5 (J48)
 - ▶ Random Forest
 - ▶ Decision Tree
- ▶ Rule Learners
 - ▶ Ripper (jRip)
- ▶ Support Vector Machines
 - ▶ SVM/LibSVM
- ▶ Bayesian Nets

Weka Makes some distinction

- ▶ Bayes
 - ▶ probabilities based on priors
- ▶ Functions
 - ▶ learn functions from data
- ▶ Lazy
 - ▶ when evaluating just go over the data again
- ▶ Meta
 - ▶ Combine classifiers together
- ▶ Misc
 - ▶ Whatever
- ▶ Rules
 - ▶ Learn boolean rules
- ▶ Trees
 - ▶ learn decision trees

Learners operate on different classes and values

- ▶ Some learners are boolean (True/False or 0/1)
- ▶ Some learners are nominal (class) (A/B/C/..)
- ▶ Some learners learn counts (1,2,3,..)
- ▶ Some learners learn real functions ($Y = b + ax$)

ZeroR Learner

- ▶ The smartest monkey
- ▶ Always chooses the class with the largest number of entities
- ▶ Good as a base line.
- ▶ You have to beat ZeroR.

C4.5/J48

- ▶ Produces a decision tree
- ▶ The model is code and interpretable
- ▶ Sometimes trees are too big.
- ▶ each branch is a conditional
- ▶ each leaf is a class

JRip/Ripper

- ▶ learns and prunes a small set of rules
- ▶ copy & paste into code

Naive Bayes

- ▶ Asks the question what is the probability of this value belonging to this class?
- ▶ multiplies all of these probabilities together

Logistic Regression

- ▶ We've already discussed this
- ▶ Regression used for true false

K-NN

- ▶ nearest neighbor
- ▶ use euclidean distance to find the

SVM

- ▶ support vector machine
- ▶ increase the dimensionality of your data to find ways to segment it in higher dimensional space
- ▶ tunable. Works well.
- ▶

Matrix of classification

- ▶ True Positives (TP) - An action or label is properly applied
 - ▶ A classifier for buggy code says buggy code is buggy
- ▶ True negative (TN) - An action or label is properly not applied
 - ▶ A classifier for buggy code says **NOT** buggy code is **NOT** buggy
- ▶ False positive (FP) - An action or label is improperly applied
 - ▶ A classifier for buggy code says NOT buggy code **IS** buggy
- ▶ False negative (FN) - An action or label is improperly NOT applied
 - ▶ A classifier for buggy code says buggy code **IS NOT** buggy

Accuracy

- ▶ Given X things how often is out automated tool right?
- ▶ E.g. given 100 samples of not working source code
 - ▶ how good is our tool at fixing the source code?
- ▶ Answer: correct / total
- ▶ $TP / (TP+TN+FN+FP)$
- ▶ $TP / \text{Everything}$
- ▶ Bad in situations where 90% of the dataset is positive
 - ▶ you just guess positive and you get 90%!
- ▶ If 90% of your data is 1 class you want better than 90% accuracy
- ▶ How many classifications were correct?
- ▶ Bad for class imbalance

Kappa

- ▶ Cohen's Kappa
- ▶ like correlation
- ▶ agreement between classifier and actual data
- ▶ Very good for class imbalance
- ▶ Check it out on Wikipedia

https://en.wikipedia.org/wiki/Cohen's_kappa

Precision

- ▶ How many of your classifications are right
- ▶ Of what was evaluated or returned what are relevant?
 - ▶ e.g. of the buggy code snippets returned how many are actually buggy?
- ▶ When I give you a positive, how right am I?
- ▶ $TP / (TP + FP)$
- ▶ Ignores the fact that I missed lots of buggy code.

Recall

- ▶ How much of the class did you find
- ▶ Might depend on the class
- ▶ You can have high precision for a class and have low recall
- ▶ Of what was evaluated or returned did I at least return most of what was relevant?
 - ▶ e.g. of the buggy code snippets returned did I return MOST of them
- ▶ Can only use when you know the population size
- ▶ When I return results do I return most of relevant results?
- ▶ $TP / (TP + FN)$

F-1 Measure

- ▶ Combination of Precision and Recall
- ▶ Geometric mean
- ▶ Can tune to one or the other
- ▶ Can I take precision and recall and balance them?
- ▶ $F1 = 2 * \text{Precision} * \text{Recall} / (\text{Precision} + \text{Recall})$
 - ▶ geometric mean of precision and recall

TP/FP Rate

- ▶ True Positives
- ▶ True Negatives
- ▶ Actual accuracy for all classes

ROC Area

- ▶ Area under the Receiver Operating Characteristic Curve
- ▶ We plot True Positive versus True Negative
- ▶ sensitivity (TPR) versus specificity (TNR)
- ▶ AUC ROC 0.5 - garbage
- ▶ AUC ROC 0.7 - good

More resources

- ▶ The wikipedia page is actually great
 - ▶ `https://en.wikipedia.org/wiki/Precision_and_recall`
 - ▶ `https://en.wikipedia.org/wiki/Cohen's_kappa`

Coffee Stain

- ▶ Find the coffee stain!
- ▶ 2 classes
- ▶ load `./data/coffee-ring/coffee-ring.arff`
- ▶ Try ZeroR
- ▶ Try 1BK
- ▶ Try SVM

Captcha

- ▶ Multiple classes (26 characters!)
- ▶ load `./data/captcha/char3.arff`
- ▶ Try ZeroR
- ▶ Try NaiveBayes
- ▶ Try SVM

Is a document reliability relevant?

- ▶ load `./data/general/pgsqla_reliability_smallerdataset.arff`
- ▶ Word based
- ▶ Can you predict if it is reliability related?
- ▶ too many features!

Dupe Bugs in Open Office

- ▶ load `./data/dupe-bugs/offswe__-.arff`
- ▶ Comparisons, can you tell which comparison will be a dupe bug or not?

How to handle text :((1/2)

- ▶ load data/triage/data/angular.js/largewo.arff
- ▶ None of this is useful!
- ▶ delete id
- ▶ Filter the owner to nominal
 - ▶ click owner
 - ▶ Weka -> Filters -> Unsupervised -> Attribute -> StringToNomial
 - ▶ click the arguments and change the index to 1
 - ▶ click apply
- ▶ ...

How to handle text :((2/2)

- ▶ Filter content to words!
 - ▶ click content
 - ▶ click filter
 - ▶ Weka -> Filters -> Unsupervised -> Attribute -> StringToWordVector
 - ▶ click arguments
 - ▶ click attributeindices
 - ▶ change to 2 or last
 - ▶ click OK
 - ▶ click apply
- ▶ Go classify and try naivebayes

ARFF Files

- ▶ Class should be the last element of the data
- ▶ Like CSV but with a type header
- ▶ String, Bool, Char, Class, Int, Float types
 - ▶ note different types for different types of jobs